

Setting up TPU for faster processing.

[illegible]

```
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:GPU:0)
Number of replicas: 8
```


Introduction 💡:

The aim of this project is to create a deep learning algorithm that is able to accurately characterize pneumonia via chest x-rays for pediatric patients. The data set that is being used is coming from Kaggle.com and is being utilized for this project.

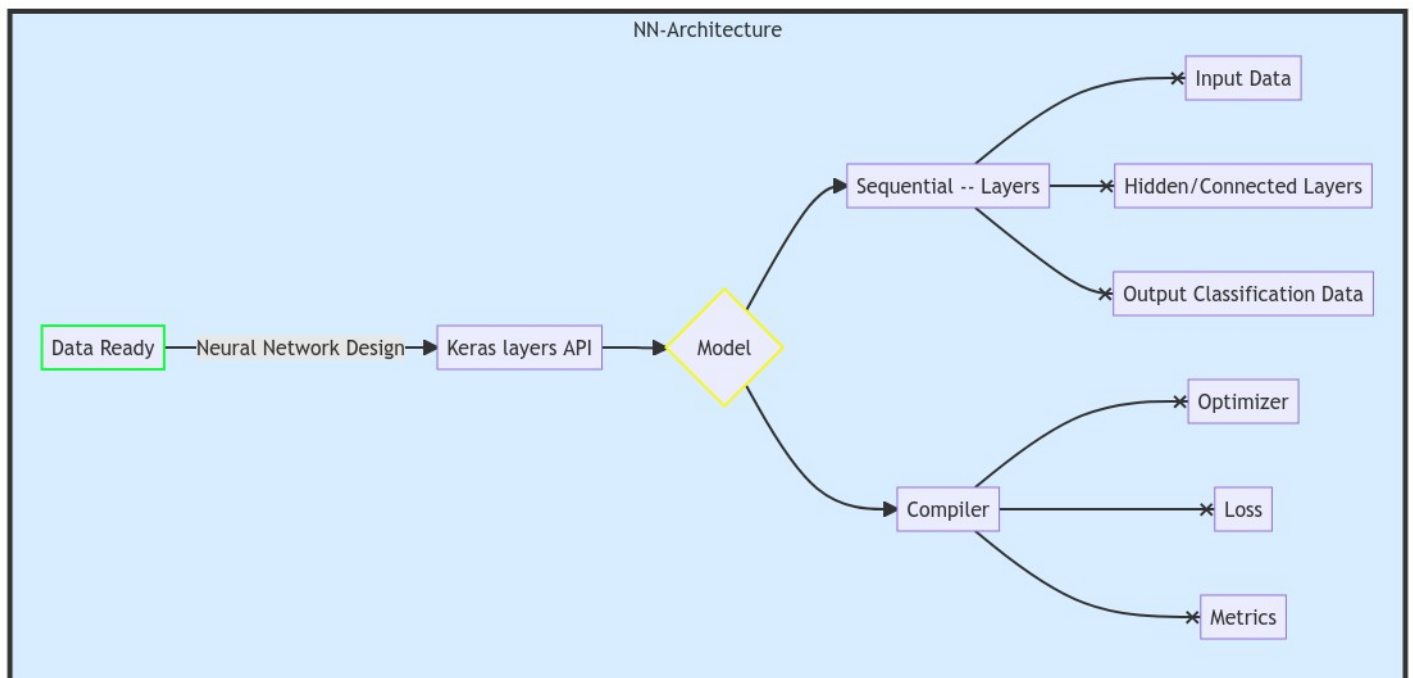
▼ Process 📅 :

Workflows illustrating the process typically taken for creating a DL project.

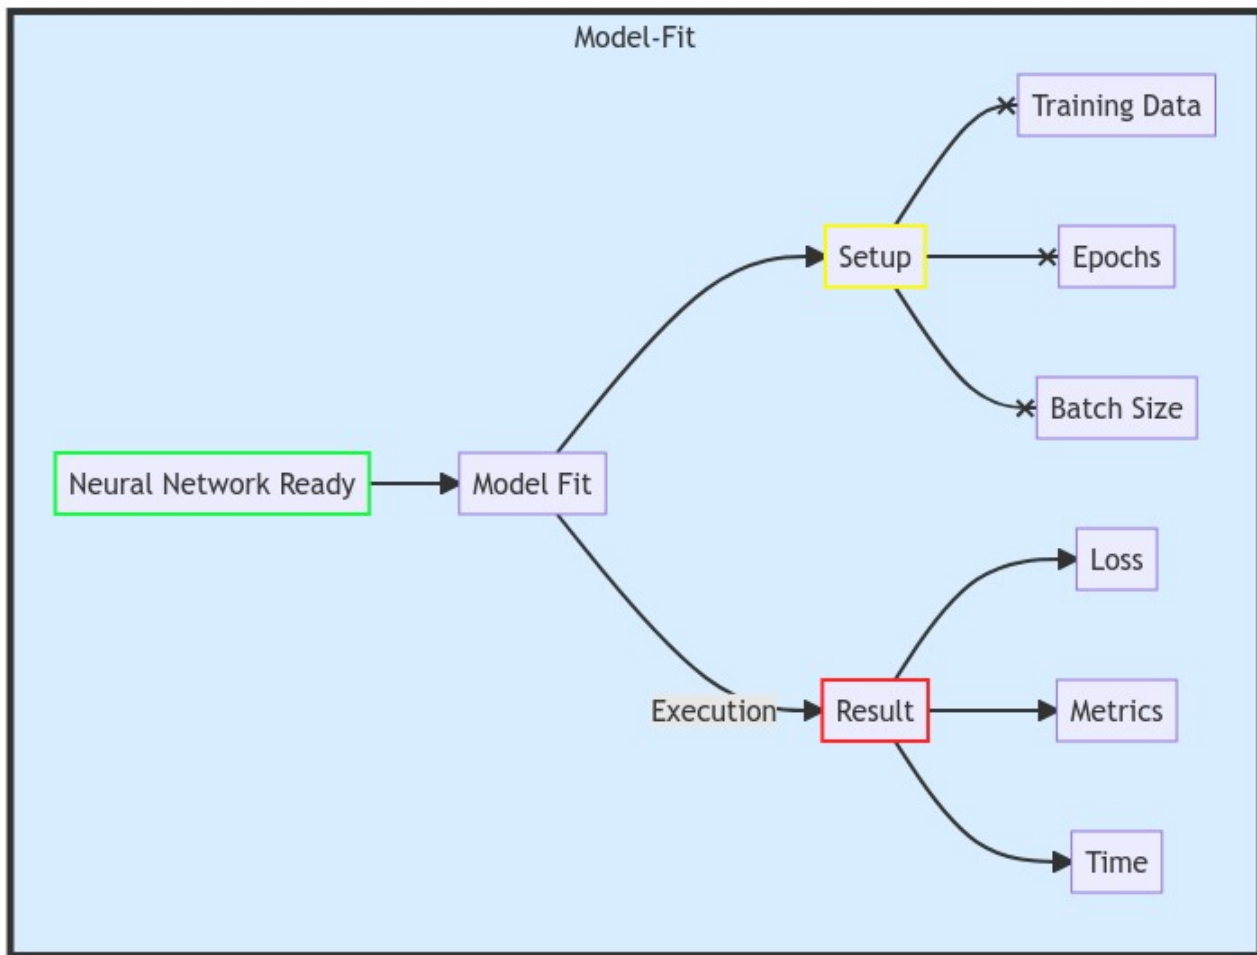
Data Preparation



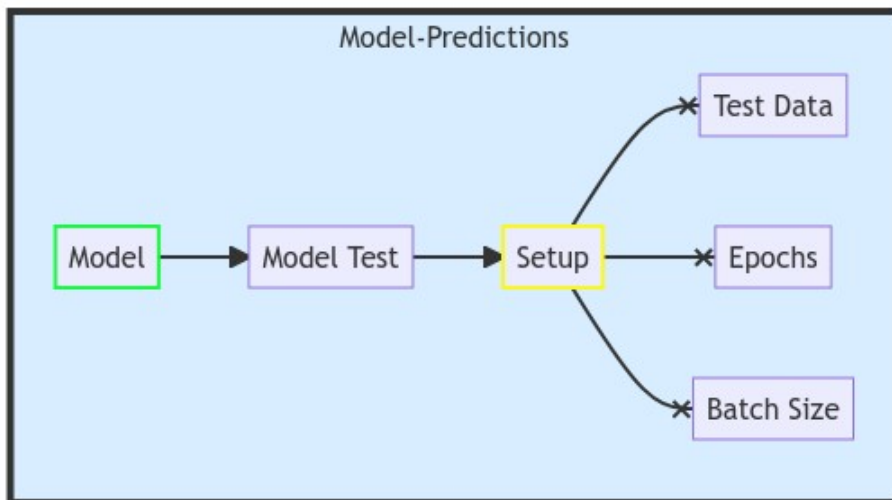
Neural Network Architecture

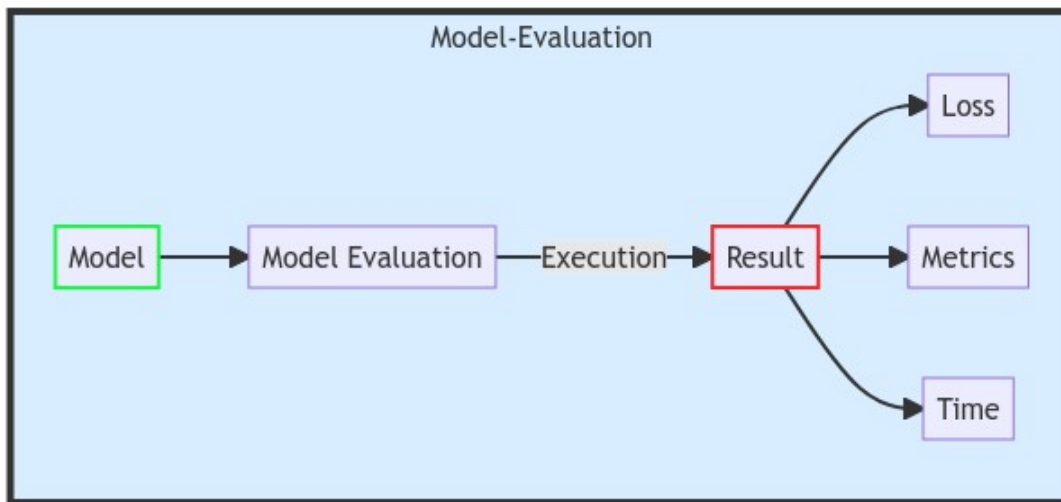


Fitting the Model



Model Predictions and Evaluation





```
!pip install colorama
!pip install opendatasets
```

Collecting colorama

Downloading colorama-0.4.4-py2.py3-none-any.whl (16 kB)

Installing collected packages: colorama

Successfully installed colorama-0.4.4

Collecting opendatasets

Downloading opendatasets-0.1.22-py3-none-any.whl (15 kB)

Requirement already satisfied: click in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: kaggle in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: certifi in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: urllib3 in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: python-slugify in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: python-dateutil in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from opendatasets)

Installing collected packages: opendatasets

Successfully installed opendatasets-0.1.22

```
import numpy as np
import pandas as pd
import re
import os
import opendatasets as od
import matplotlib.pyplot as plt
from tensorflow import keras
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from sklearn.model_selection import train_test_split
import colorama
```

▼ Execution 🤖

```
#Importing data set kaggle API for initial data review
dataset = 'https://www.kaggle.com/datasets/paultimothymooney/chest-xray-pneumonia'
od.download(dataset)

Please provide your Kaggle credentials to download this dataset. Learn more: http://bit.ly
Your Kaggle username: mdespinoza
Your Kaggle Key: .....
Downloading chest-xray-pneumonia.zip to ./chest-xray-pneumonia
100%|██████████| 2.29G/2.29G [00:12<00:00, 200MB/s]
```

▼ Data Review

In this section we will perform an assesment of the dataset without manipulating any of the data.

After importing the data we will need to go through and identify the following:

- How many data sets are we given?
- How many classes are there in the dataset?

```
data_dir = './chest-xray-pneumonia'
os.listdir(data_dir)

['chest_xray']

# Identifying the number folders
os.listdir('./chest-xray-pneumonia/chest_xray')

['val', 'train', 'test', '__MACOSX', 'chest_xray']

# Identifying the categories
os.listdir('./chest-xray-pneumonia/chest_xray/test')

['NORMAL', 'PNEUMONIA']

CATEGORIES = ["PNEUMONIA", "NORMAL"]

# Identify the number of images per folder
test_set = ('./chest-xray-pneumonia/chest_xray/test')
train_set = ('./chest-xray-pneumonia/chest_xray/train')
val_set = ('./chest-xray-pneumonia/chest_xray/val')

# function to calculate the number of images in a dataset folder
```

```

def count_xray_images(title, dataset, color):
    print(color + title)
    pneumonia = len(os.listdir(os.path.join(dataset, 'PNEUMONIA')))
    normal = len(os.listdir(os.path.join(dataset, 'NORMAL')))
    print(f"Pneumonia = {pneumonia}")
    print(f"Percent Pneumonia = {pneumonia/(pneumonia + normal)*100:.1f}%")
    print(f"Normal = {normal}")
    print(f"Percent Normal = {normal/(pneumonia + normal)*100:.1f}%")
    print(f"TOTAL = {pneumonia + normal}")

# function to create a bargraph for the number of images per category in any dataset folder
def data_visual(title, dataset, y_limit):
    pneumonia = len(os.listdir(os.path.join(dataset, 'PNEUMONIA')))
    normal = len(os.listdir(os.path.join(dataset, 'NORMAL')))
    x_pos = CATEGORIES
    y = (normal, pneumonia)
    plt.title(title)
    # define color for bars
    c = ("#FF0000", "#ADD8E6")
    # Adding labels, and the title
    plt.xlabel("Category")
    plt.ylabel("Num. of X-Ray Images")
    plt.bar(x_pos, y, color = c, ec = "black")
    plt.ylim(y_limit)

    for i in range(len(x_pos)):
        plt.text(i, y[i], y[i], ha="center", va="bottom")

    plt.show()


# Calling the function to provide the calculations for each folder and visualization
count_xray_images("Test Folder", test_set, colorama.Fore.BLUE)
data_visual("Test Folder", test_set,[0, 450])
count_xray_images("Train Folder", train_set, colorama.Fore.GREEN)
data_visual("Train Folder", train_set,[400, 4500])
count_xray_images("Val Folder", val_set, colorama.Fore.RED)
data_visual("Val Folder", val_set,[0, 10])

```

```

Test Folder
Pneumonia = 390
Percent Pneumonia = 62.5%
Normal = 234
Percent Normal = 37.5%
TOTAL = 624


```



```

Train Folder
Pneumonia = 3875
Percent Pneumonia = 74.3%
Normal = 1341
Percent Normal = 25.7%
TOTAL = 5216

```



```

Val Folder
Pneumonia = 8
Percent Pneumonia = 50.0%

```

```

# function to display images
def image_xray_batch(path1, path2, title):
    image_list = os.listdir(path1)
    image = path2

    plt.figure(figsize=(5,5))

    for i in range(16):
        plt.subplot(4, 4, i + 1)
        img = plt.imread(os.path.join(image, image_list[i]))
        plt.imshow(img, cmap='gray')
        plt.axis('off')

```


```

plt.title(title)
plt.tight_layout()

image_xray_batch(
'./chest-xray-pneumonia/chest_xray/train/NORMAL',
'./chest-xray-pneumonia/chest_xray/train/NORMAL',
"PNEUMONIA")

image_xray_batch(
'./chest-xray-pneumonia/chest_xray/train/PNEUMONIA',
'./chest-xray-pneumonia/chest_xray/train/PNEUMONIA',
"NORMAL")

```



```

norm_img = os.listdir("./chest-xray-pneumonia/chest_xray/train/NORMAL")[0]
norm_dir = "./chest-xray-pneumonia/chest_xray/train/NORMAL"
sample_img = plt.imread(os.path.join(norm_dir, norm_img))
plt.imshow(sample_img, cmap='gray')
plt.colorbar()


```



```
plt.title('X Ray Image')
```

```
print(f"The max px value = {sample_img.max():.1f} and the min = {sample_img.min():.1f}")
```

The max px value = 255.0 and the min = 0.0



Data Review Summary:

The scope of this project is based on a classification problem that is to assess and predict if a Pediatric patient has pneumonia or not (2 Categories = Pneumonia, Normal). However in reviewing the training data it is evident that the observations for "Normal" (majority class), is significantly higher than the "Pneumonia" (minority class) observations. Which implies that the dataset is imbalanced and will need to be corrected in the data preparation to avoid any bias towards the "Normal" class.

The process to balance the data between both classes would be to adjust the weights of each class.

Also, the validation data is very small $n=16$ and not sufficient for creating a model that we could trust. Thus in the next section we will also split the training and validation data using a standard split of 80:20 to create a larger validation dataset

▼ Data Preparation 12 34

```
#accessing the data via googlecloudstorage - which allows for the utilization of Google's TPU
```

```
GCS_PATH = 'gs://kds-7d8d6352305c640b4084285616169079853d84b5df55d6be0bc51958'
```

```
# paths defined to variables
```

```
test_image_folder = (GCS_PATH + '/chest_xray/test/**')
```

```
train_set_folder = (GCS_PATH + '/chest_xray/train/**')
```

```
val_set_folder = (GCS_PATH + '/chest_xray/val/**')
```

```
# Split our training dataset into a training and validation set via "train_test_split" by a spl
```

```
train_files = tf.io.gfile.glob(str(train_set_folder))
```

```
train_files.extend(tf.io.gfile.glob(str(val_set_folder)))
```

```

training, validation = train_test_split(train_files, test_size=0.2)
# Verification of (n) of train dataset in each category folder created after the split.
normal_n = len([train_files for train_files in training if "NORMAL" in train_files])
print("Normal images: " + str(normal_n))

pneu_n = len([train_files for train_files in training if "PNEUMONIA" in train_files])
print("PNEUMONIA images: " + str(pneu_n))

    Normal images: 1063
    PNEUMONIA images: 3122

# creating slices from our defined dataset arrays

# training
data_struct_train = tf.data.Dataset.from_tensor_slices(training)
# validation
data_struct_val = tf.data.Dataset.from_tensor_slices(validation)
# test
test_data_struct = tf.data.Dataset.list_files(str(test_image_folder))

# train data
train_count = tf.data.experimental.cardinality(data_struct_train).numpy()
print("Training (80%) = " + str(train_count))
# val data
val_count = tf.data.experimental.cardinality(data_struct_val).numpy()
print("Validating (20%) = " + str(val_count))
# test data
test_count = tf.data.experimental.cardinality(test_data_struct).numpy()
print("Test = " + str(test_count ))

    Training (80%) = 4185
    Validating (20%) = 1047
    Test = 624

# defining data labels as an array
categories_labels = np.array([str(tf.strings.split(item, os.path.sep)[-1].numpy())[2:-1]
                             for item in tf.io.gfile.glob(str(GCS_PATH + '/chest_xray/train/*')

categories_labels

    array(['NORMAL', 'PNEUMONIA'], dtype='<U9')

# Setting dataset to be labeled with categories(0)|(1)
def define_label(file_path):
    parts = tf.strings.split(file_path, os.path.sep)
    return parts[-2] == "PNEUMONIA"

# defining image size constant
IMAGE_SIZE = [180, 180]

# Establish uniformity with the images (3D Tensor) and scaling them down
# to smaller sizes.

```

```

def img_conversion(img):
    img = tf.image.decode_jpeg(img, channels=3)
    img = tf.image.convert_image_dtype(img, tf.float32)
    return tf.image.resize(img, IMAGE_SIZE)

def path(file_path):
    label = define_label(file_path)
    img = tf.io.read_file(file_path)
    img = img_conversion(img)
    return img, label

autotune = tf.data.experimental.AUTOTUNE

BATCH_SIZE = 16 * tpu_strategy.num_replicas_in_sync

train_data_st = data_struct_train.map(path, num_parallel_calls=autotune)

val_data_st = data_struct_train.map(path, num_parallel_calls=autotune)

test_data_st = test_data_struct.map(path, num_parallel_calls=autotune)

test_data_st = test_data_st.batch(BATCH_SIZE)

# verify data normalization = image shape and label
for image, label in train_data_st.take(1):
    print("Image shape: ", image.numpy().shape)
    print("Label: ", label.numpy())

    Image shape:  (180, 180, 3)
    Label:  True

def prep_data_pre_training(ds, cache=True, shuffle_buffer_size=1000):
    # This is a small dataset, only load it once, and keep it in memory.
    # use `.cache(filename)` to cache preprocessing work for datasets that don't
    # fit in memory.
    if cache:
        if isinstance(cache, str):
            ds = ds.cache(cache)
        else:
            ds = ds.cache()

    ds = ds.shuffle(buffer_size=shuffle_buffer_size)

    # Repeat forever
    ds = ds.repeat()

    ds = ds.batch(BATCH_SIZE)

    # `prefetch` lets the dataset fetch batches in the
    # background while the model is training.
    ds = ds.prefetch(buffer_size=autotune)

```

```

train_data_st = prep_data_pre_training(train_data_st)
val_data_st = prep_data_pre_training(val_data_st)

image_batch, label_batch = next(iter(train_data_st))

# Adjusting the weights for both classes to compensate for the
# data not being equal
weight_for_0 = (1 / normal_n)*(train_count)/2.0
weight_for_1 = (1 / pneu_n)*(train_count)/2.0

class_weight = {0: weight_for_0, 1: weight_for_1}

print('Normal Weight: {:.2f}'.format(weight_for_0))
print('Pneumonia Weight: {:.2f}'.format(weight_for_1))

Normal Weight: 1.97
Pneumonia Weight: 0.67

```

▼ Neural Network Architecture 🧠

```

def build_model():
    # The model is going to consist of a sequence of layers as
    # noted by tf.keras.Sequential
    model = tf.keras.Sequential([
        tf.keras.Input(shape=(IMAGE_SIZE[0], IMAGE_SIZE[1], 3)), # expected input shape

        # Conv2D = convolution layer
        # setting filters (16) and kernel size (3)
        # padding defined as 'same' which implies the preservation
        # of spatial dimensions
        # with activation Rectified Linear Unit = 0 or greater is the output
        tf.keras.layers.Conv2D(16, 3, activation='relu', padding='same'),
        tf.keras.layers.Conv2D(16, 3, activation='relu', padding='same'),
        # Maxpool2D - pooling of max values
        tf.keras.layers.MaxPool2D(),

        tf.keras.Sequential([
            tf.keras.layers.SeparableConv2D(32, 3, activation='relu', padding='same'),
            tf.keras.layers.SeparableConv2D(32, 3, activation='relu', padding='same'),
            tf.keras.layers.BatchNormalization(),
            tf.keras.layers.MaxPool2D()]),

        tf.keras.Sequential([
            tf.keras.layers.SeparableConv2D(64, 3, activation='relu', padding='same'),
            tf.keras.layers.SeparableConv2D(64, 3, activation='relu', padding='same'),
            tf.keras.layers.BatchNormalization(),
            tf.keras.layers.MaxPool2D()]),

        tf.keras.Sequential([
            tf.keras.layers.SeparableConv2D(128, 3, activation='relu', padding='same'),
            tf.keras.layers.SeparableConv2D(128, 3, activation='relu', padding='same'),

```

```

tf.keras.layers.BatchNormalization(),
tf.keras.layers.MaxPool2D()]),

tf.keras.layers.Dropout(0.2),

tf.keras.Sequential([
tf.keras.layers.SeparableConv2D(256, 3, activation='relu', padding='same'),
tf.keras.layers.SeparableConv2D(256, 3, activation='relu', padding='same'),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.MaxPool2D()]),

# prevents over fitting the model
tf.keras.layers.Dropout(0.2),

# reshape the input
tf.keras.layers.Flatten(),

tf.keras.Sequential([
# Dense implies adding a layer of neurons
tf.keras.layers.Dense(512, activation='relu'),
# normalizing the data
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Dropout(0.7)]),

tf.keras.Sequential([
tf.keras.layers.Dense(128, activation='relu'),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Dropout(0.5)]),

tf.keras.Sequential([
tf.keras.layers.Dense(64, activation='relu'),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Dropout(0.3)]),

tf.keras.layers.Dense(1, activation='sigmoid')
])

return model

# Distribute training across the TPU
with tpu_strategy.scope():
    model = build_model()
    # defining metrics to be as outputs in the neural network
    METRICS = ['accuracy', # fration of predictions the model gets correct
               # measure what proportion of "+" id's were correct
               tf.keras.metrics.Precision(name='precision'),
               # measure what prportion of actual positives weere id correctly
               tf.keras.metrics.Recall(name='recall')]

#setting up compiler
model.compile(
    # Applying gradient descent algorithm
    optimizer='adam',
    # cross entropy loss for binary (0 or 1 )
    # classifications = NORMAL or PNEUMONA

```

```
loss='binary_crossentropy',
```

▼ Fitting the Model

```
EPOCHS = 15
```

```
#initial - 1st model fit training
```

```
history = model.fit(  
    train_data_st,  
    steps_per_epoch= train_count // BATCH_SIZE, # num of sample / batch size  
    epochs=EPOCHS,  
    validation_data=val_data_st,  
    validation_steps= val_count // BATCH_SIZE, # num of sample / batch size  
    class_weight=class_weight, # data imbalance correction by weights  
)
```

```
Epoch 1/15  
32/32 [=====] - 145s 3s/step - loss: 0.5094 - accuracy: 0.7363 -  
Epoch 2/15  
32/32 [=====] - 45s 1s/step - loss: 0.3118 - accuracy: 0.8652 - r  
Epoch 3/15  
32/32 [=====] - 46s 1s/step - loss: 0.2411 - accuracy: 0.9043 - r  
Epoch 4/15  
32/32 [=====] - 46s 1s/step - loss: 0.2158 - accuracy: 0.9163 - r  
Epoch 5/15  
32/32 [=====] - 46s 1s/step - loss: 0.1961 - accuracy: 0.9294 - r  
Epoch 6/15  
32/32 [=====] - 47s 2s/step - loss: 0.1702 - accuracy: 0.9399 - r  
Epoch 7/15  
32/32 [=====] - 49s 2s/step - loss: 0.1733 - accuracy: 0.9333 - r  
Epoch 8/15  
32/32 [=====] - 47s 2s/step - loss: 0.1908 - accuracy: 0.9355 - r  
Epoch 9/15  
32/32 [=====] - 47s 2s/step - loss: 0.1680 - accuracy: 0.9414 - r  
Epoch 10/15  
32/32 [=====] - 52s 2s/step - loss: 0.1384 - accuracy: 0.9451 - r  
Epoch 11/15  
32/32 [=====] - 46s 1s/step - loss: 0.1397 - accuracy: 0.9500 - r  
Epoch 12/15  
32/32 [=====] - 48s 2s/step - loss: 0.1272 - accuracy: 0.9556 - r  
Epoch 13/15  
32/32 [=====] - 48s 2s/step - loss: 0.1313 - accuracy: 0.9536 - r  
Epoch 14/15  
32/32 [=====] - 47s 2s/step - loss: 0.1264 - accuracy: 0.9556 - r  
Epoch 15/15  
32/32 [=====] - 49s 2s/step - loss: 0.1105 - accuracy: 0.9607 - r
```

```
model.summary()
```

```
Model: "sequential_7"
```

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 180, 180, 16)	448
conv2d_1 (Conv2D)	(None, 180, 180, 16)	2320

max_pooling2d (MaxPooling2D	(None, 90, 90, 16)	0
)		
sequential (Sequential)	(None, 45, 45, 32)	2160
sequential_1 (Sequential)	(None, 22, 22, 64)	7392
sequential_2 (Sequential)	(None, 11, 11, 128)	27072
dropout (Dropout)	(None, 11, 11, 128)	0
sequential_3 (Sequential)	(None, 5, 5, 256)	103296
dropout_1 (Dropout)	(None, 5, 5, 256)	0
flatten (Flatten)	(None, 6400)	0
sequential_4 (Sequential)	(None, 512)	3279360
sequential_5 (Sequential)	(None, 128)	66176
sequential_6 (Sequential)	(None, 64)	8512
dense_3 (Dense)	(None, 1)	65

```

=====
Total params: 3,496,801
Trainable params: 3,494,433
Non-trainable params: 2,368

```

```

# Save keras model so that model may be used later to continue training the data
# from the saved state.

```

```

checkp_cb = tf.keras.callbacks.ModelCheckpoint("pneumonia_model.h5",
                                              save_best_only=True)

```

```

# training will terminate if loss does not improve after "5" epochs
early_stop_cb = tf.keras.callbacks.EarlyStopping(patience=5,
                                                restore_best_weights=True)

```

```

# defining learning rate decay to be applied accross the 2nd model fit training
# until local minima is obtained

```

```

def exponential_decay(lr0, s):
    def exponential_decay_fn(epoch):
        return lr0 * 0.1 ** (epoch / s)
    return exponential_decay_fn

```

```

exponential_decay_fn = exponential_decay(0.01, 20)

```

```

lr_scheduler = tf.keras.callbacks.LearningRateScheduler(exponential_decay_fn)

```

```

#second - 2nd model fit training

```

```

history = model.fit(
    train_data_st,
    steps_per_epoch= train_count // BATCH_SIZE,
    # increased batch size as callbacks has been added to stop training after
    # defined criteria has been achieved.
    epochs= 100,

```

```
validation_data=val_data_st,  
validation_steps= val_count // BATCH_SIZE,  
class_weight=class_weight,  
callbacks=[checkp_cb, early_stop_cb, lr_scheduler]
```

)

```
Epoch 1/100  
32/32 [=====] - 51s 2s/step - loss: 0.2618 - accuracy: 0.8999 - f  
Epoch 2/100  
32/32 [=====] - 49s 2s/step - loss: 0.2102 - accuracy: 0.9170 - f  
Epoch 3/100  
32/32 [=====] - 49s 2s/step - loss: 0.1995 - accuracy: 0.9175 - f  
Epoch 4/100  
32/32 [=====] - 51s 2s/step - loss: 0.2140 - accuracy: 0.9146 - f  
Epoch 5/100  
32/32 [=====] - 49s 2s/step - loss: 0.1648 - accuracy: 0.9333 - f  
Epoch 6/100  
32/32 [=====] - 48s 2s/step - loss: 0.1520 - accuracy: 0.9375 - f  
Epoch 7/100  
32/32 [=====] - 50s 2s/step - loss: 0.1162 - accuracy: 0.9529 - f  
Epoch 8/100  
32/32 [=====] - 49s 2s/step - loss: 0.1115 - accuracy: 0.9570 - f  
Epoch 9/100  
32/32 [=====] - 50s 2s/step - loss: 0.1197 - accuracy: 0.9497 - f  
Epoch 10/100  
32/32 [=====] - 47s 2s/step - loss: 0.1044 - accuracy: 0.9602 - f  
Epoch 11/100  
32/32 [=====] - 50s 2s/step - loss: 0.1055 - accuracy: 0.9556 - f  
Epoch 12/100  
32/32 [=====] - 49s 2s/step - loss: 0.0914 - accuracy: 0.9683 - f  
Epoch 13/100  
32/32 [=====] - 51s 2s/step - loss: 0.0852 - accuracy: 0.9634 - f  
Epoch 14/100  
32/32 [=====] - 50s 2s/step - loss: 0.0898 - accuracy: 0.9719 - f  
Epoch 15/100  
32/32 [=====] - 50s 2s/step - loss: 0.0907 - accuracy: 0.9692 - f  
Epoch 16/100  
32/32 [=====] - 52s 2s/step - loss: 0.0650 - accuracy: 0.9768 - f  
Epoch 17/100  
32/32 [=====] - 50s 2s/step - loss: 0.0558 - accuracy: 0.9788 - f  
Epoch 18/100  
32/32 [=====] - 49s 2s/step - loss: 0.0826 - accuracy: 0.9695 - f  
Epoch 19/100  
32/32 [=====] - 51s 2s/step - loss: 0.0546 - accuracy: 0.9814 - f  
Epoch 20/100  
32/32 [=====] - 51s 2s/step - loss: 0.0687 - accuracy: 0.9775 - f  
Epoch 21/100  
32/32 [=====] - 50s 2s/step - loss: 0.0505 - accuracy: 0.9812 - f  
Epoch 22/100  
32/32 [=====] - 51s 2s/step - loss: 0.0446 - accuracy: 0.9834 - f  
Epoch 23/100  
32/32 [=====] - 49s 2s/step - loss: 0.0397 - accuracy: 0.9866 - f  
Epoch 24/100  
32/32 [=====] - 49s 2s/step - loss: 0.0521 - accuracy: 0.9839 - f  
Epoch 25/100  
32/32 [=====] - 51s 2s/step - loss: 0.0432 - accuracy: 0.9814 - f  
Epoch 26/100  
32/32 [=====] - 50s 2s/step - loss: 0.0367 - accuracy: 0.9893 - f  
Epoch 27/100  
32/32 [=====] - 51s 2s/step - loss: 0.0364 - accuracy: 0.9893 - f  
Epoch 28/100  
32/32 [=====] - 48s 2s/step - loss: 0.0293 - accuracy: 0.9885 - f
```


Epoch 29/100

32/32 [=====] - 50s 2s/step - loss: 0.0278 - accuracy: 0.9885 - r

```
model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 180, 180, 16)	448
conv2d_1 (Conv2D)	(None, 180, 180, 16)	2320
max_pooling2d (MaxPooling2D)	(None, 90, 90, 16)	0
sequential (Sequential)	(None, 45, 45, 32)	2160
sequential_1 (Sequential)	(None, 22, 22, 64)	7392
sequential_2 (Sequential)	(None, 11, 11, 128)	27072
dropout (Dropout)	(None, 11, 11, 128)	0
sequential_3 (Sequential)	(None, 5, 5, 256)	103296
dropout_1 (Dropout)	(None, 5, 5, 256)	0
flatten (Flatten)	(None, 6400)	0
sequential_4 (Sequential)	(None, 512)	3279360
sequential_5 (Sequential)	(None, 128)	66176
sequential_6 (Sequential)	(None, 64)	8512
dense_3 (Dense)	(None, 1)	65
Total params: 3,496,801		
Trainable params: 3,494,433		
Non-trainable params: 2,368		

```
# plotting training results
```

```
fig, ax = plt.subplots(1, 3, figsize=(20, 5))
```

```
ax = ax.ravel()
```

```
for i, met in enumerate(['precision', 'recall', 'accuracy',]):
```

```
    ax[i].plot(history.history[met])
```


```
    ax[i].plot(history.history['val_' + met])
```

```
    ax[i].set_title('Model {}'.format(met))
```

```
    ax[i].set_xlabel('epochs')
```

```
    ax[i].set_ylabel(met)
```

```
    ax[i].legend(['train', 'val'])
```



▼ Model Predictions

```
loss, acc, prec, rec = model.evaluate(test_data_st)
```

```
5/5 [=====] - 20s 4s/step - loss: 1.0036 - accuracy: 0.8029 - pre
```

Summary

1. The model overall has an 98.95% accuracy and a loss of 3.55%.
 - model may be improved by adjusting the CNN (Fine-tuning)
2. The model prediction accuracy is 80.29% which is less than both the training and validation accuracy. (Can be improved!)
3. The model prediction Recall is larger than the precision. Precision may be improved by adding more data that is of high quality.
 - 76.33 % of total class outcomes was predicted correctly
4. To improve the results there are two options:
 - Fine-tuning of the CNN and re-training it.
 - Capture more data

Sources

Acknowledgements Data: <https://data.mendeley.com/datasets/rscbjbr9sj/2>

License: CC BY 4.0

Citation: [http://www.cell.com/cell/fulltext/S0092-8674\(18\)30154-5](http://www.cell.com/cell/fulltext/S0092-8674(18)30154-5)

Kaggle Project: <https://www.kaggle.com/datasets/paultimothymooney/chest-xray-pneumonia>

Other Kaggle opensource review on dataset:

1. <https://www.kaggle.com/code/aakashnain/beating-everything-with-depthwise-convolution>
2. <https://www.kaggle.com/code/amyjang/tensorflow-pneumonia-classification-on-x-rays>
3. <https://towardsdatascience.com/deep-learning-for-detecting-pneumonia-from-x-ray-images-fc9a3d9fdb8>
4. <https://towardsdatascience.com/detecting-covid-19-induced-pneumonia-from-chest-x-rays-with-transfer-learning-an-implementation-311484e6afc1>

Colab TPU Use With Flowes modeling:

<https://colab.research.google.com/notebooks/tpu.ipynb#scrollTo=LtAVr-4CP1rp>

Webiste:

https://colab.research.google.com/notebooks/snippets/accessing_files.ipynb#scrollTo=z1_FuDjAozF1